

Point-in-Time Multi-Agent LLM Trading: A Reproducible LangGraph Stack with Regime-Conditioned Case Study Evaluation and Fair Backtesting

Rohit Edward Martires

Independent researcher, Goa, India

`rohitmartires14@gmail.com`

29 May 2026

Abstract

Published LLM trading stacks are often hard to reproduce, rarely enforce temporal data boundaries, and seldom report net returns after fees. We present **Hindsight 20/20**, an open **LangGraph** stack with structured **Pydantic** outputs, point-in-time data access, optional **ticker anonymization**, and a fee-aware **paper backtest**. The primary contribution is reproducible systems engineering with frozen CSV artifacts—not a universe-scale alpha benchmark. As an illustrative case study, we evaluate the **full** pipeline on **RELIANCE.NS** and **TCS.NS** in **four pre-specified regime windows** (two tickers $\times$ exogenous bear/bull labels; **one trajectory per window**, no seed averaging). In both bear windows, the agent loses less than buy-and-hold with lower drawdown; bull outcomes diverge (RELIANCE lags B&H; TCS nearly tracks it). We interpret heterogeneous bull results as an architectural boundary of defensive

multi-stage coordination. Code: <https://github.com/RMartires/hindsight>

Keywords: algorithmic trading; large language models; multi-agent systems; reproducibility; point-in-time data; market regimes.

0.1 Introduction

Researchers are experimenting with LLMs inside trading workflows amid rapid growth in LLM-driven agent systems (Bădică et al., 2025), but rarely report how the **same stack** behaves across **different underlying trends**, enforce point-in-time data at the adapter layer, or report **net** returns after stated fees.

What this paper is. **Hindsight 20/20** extends **TradingAgents** (Xiao et al., 2024) with schema-constrained outputs, a simulation end date, optional ticker anonymization, and a scriptable fee-aware backtest tied to **frozen CSV artifacts**.

What this paper is not. We do **not** claim state-of-the-art risk-adjusted returns, regime prediction, or generalization from a broad universe. Each regime window is **N=1** (no multi-seed variance). Frontier benchmarks (Jeon and Lee, 2026; Miyazaki et al., 2026) average many seeds across index universes; our scope is narrower and labeled as such (Section 4.1).

We evaluate the **full** pipeline on **RELIANCE.NS** and **TCS.NS** in **four pre-specified regime windows** (bear and bull per ticker), each from \$100,000 paper cash.

0.2 Related Work

LLM agents and classic MAS. Multi-agent trading frameworks orchestrate specialist roles (Xiao et al., 2024; Yu et al., 2023; Yang et al., 2023). Bădică et al. (2025) contrast emergent LLM agents with classic BDI multi-agent systems. **TradingAgents** supplies our graph skeleton; **AlphaAgents** (Zhao et al., 2025) validates institution-style role separation; **QuantAgents** (Li et al., 2025) emphasizes manager-led meetings—comparators for our

debate loop, but neither foregrounds adapter-level point-in-time routing or frozen schedule CSVs.

Methodological standards and bias control. Nguyen and Pham (2026) propose Consolidated Minimum Standards for LLM financial multi-agent evaluation—contamination control, net-of-cost returns, and regime awareness—and introduce the Coordination Breakeven Spread (CBS). Bailey et al. (2014) warn that undisclosed configuration search makes strong backtests easy to overfit; we respond with pre-specified windows and frozen artifacts. While TradingAgents provides the foundational graph, Hindsight 20/20 implements the point-in-time policy and net-cost ledger those standards recommend.

Anonymization and fine-grained tasks. Jeon and Lee (2026) (*BlindTrade*) show that identifiers such as AAPL can trigger memorized pretraining rather than active reasoning; they anonymize S&P~500 names and report Sharpe $1.40, \pm 0.22$ across 20 seeds. Hindsight’s optional STOCK_#### aliases pursue the same goal in a LangGraph replication stack. Miyazaki et al. (2026) prove fine-grained analyst sub-tasks outperform coarse prompts; our analysts are tool-bound to indicators and statements (Section~3.2). Dewansh (2026) defines *explainable stability* and rationale persistence; our Pydantic stage outputs support audit trails from outlook to final_signal.

Benchmarks and simulation realism. Qian et al. (2026) (*Agent Market Arena*) show agent architecture often matters more than LLM backbone—motivating our fixed-model, regime-varying case study. Papadakis et al. (2025) (*StockSim*) define order-level simulation with latency and slippage; our candlestick ledger with stated fees is deliberately simpler, with StockSim-class realism as future work.

Scope contrast. Jeon and Lee (2026), Li et al. (2025), and Miyazaki et al. (2026) stress multi-seed index evaluation; Hindsight stresses reproducible temporal integrity on four pre-specified windows (Table 2).

0.3 Method / System

TradingAgentsGraph: analysts $\rightarrow$ bull/bear debate $\rightarrow$ trader $\rightarrow$ risk. Section 4 uses `PAPER_ABLATION=full. backtest_mvp.py --dates-csv` produces wide schedule CSVs.

0.3.1 Analyst task decomposition

Stage	Tools	Tasks
Market	<code>get_stock_data</code> , <code>get_indicators</code>	OHLCV + up to 8 Tier-1 indicators (SMA, EMA, MACD, RSI, Bollinger, ATR, VWMA, MFI)
Fundamentals	<code>get_fundamentals</code> , statements	Revenue, margins, balance-sheet trends
News / Sentiment	<code>get_news</code> , <code>get_global_news</code>	Company and macro news synthesis
Debate $\rightarrow$ Risk	—	Structured stances $\rightarrow$ BUY/HOLD/SELL

0.3.2 Ticker anonymization

Deterministic `STOCK_####` aliases for LLM I/O; vendor round-trip de-anonymization; news proper-noun scrubbing. Frozen E1: `enable_anonymization=True`.

0.3.3 Backtest metrics

We emphasize **net return**, **MDD**, **fee drag**, and **alpha vs buy-and-hold**. **Excess Sharpe** uses daily equity from frozen CSVs with a **6.5% p.a.** risk-free rate (India short-rate proxy)¹. Sortino and Calmar are secondary on short windows (67–111 days).

¹Annualization: $\sqrt{252}$; NSE $\approx$ 240–250 trading days.

0.4 Case Study Evaluation

0.4.1 Scope and limits (Section 4.1)

Two liquid NSE large-caps; **four regime windows**; **N=1 trajectory per window**; exogenous bear/bull labels ($\pm 10\%$ underlying B&H); pre-specified calendars in **regime-timeranges**. Not a hypothesis test about expected alpha.

0.4.2 Positioning (Table 2)

	Hindsight	BlindTrade	QuantAgents	Expert Teams
Primary goal	PIT reproducibility	Anonymized policy	Meetings	Hierarchy
Seeds	N=1/window	20	Multi-year	50
Universe	2 NSE names	S&P 500	Simulation	TOPIX 100
Costs	Net (Zerodha)	Varies	Varies	Varies

Table 2: Qualitative comparison; Hindsight is not comparable on seeds/universe.

0.4.3 Results

Table 1. Primary metrics from frozen CSVs (\$100,000 start).

Ticker	Regime	Days	B&H	Agent	Gross	Fee	α	Ag. MDD	B&H MDD
RELIANCE	Bear	111	-17.4%	-7.5%	-6.5%	1.0pp	+9.9pp	5.1%	21.0%
RELIANCE	Bull	108	+15.1%	+5.6%	+6.2%	0.6pp	-9.5pp	2.8%	11.0%
TCS	Bear	88	-20.4%	-4.8%	-2.9%	1.9pp	+15.6pp	6.8%	23.9%
TCS	Bull	67	+21.9%	+19.6%	+21.8%	2.2pp	-2.3pp	1.7%	5.5%

Table 3: Net return, gross return, fee drag, alpha vs B&H, and MDD.

Table 3. Excess Sharpe (6.5% p.a. R_f), Sortino, Calmar (post-hoc on frozen equity).

Ticker	Regime	Excess Sh. (agent)	Excess Sh. (B&H)	Sortino	Calmar
RELIANCE	Bear	-4.27	-2.51	-3.10	-0.96
RELIANCE	Bull	+0.74	+1.25	+3.67	+1.94
TCS	Bear	-2.69	-3.13	-2.25	-0.71
TCS	Bull	+3.86	+3.27	+24.22	+11.88

Table 4: Secondary risk-adjusted metrics; short windows make ratios noisy.

Patterns. Bear (2/2): smaller loss and lower MDD than B&H; SELL-heavy signals. Bull: TCS nearly tracks B&H; RELIANCE lags despite positive return (+5.6%)—fee drag (\$ \$0.6pp) does not explain the \$ \$9.5pp alpha gap; defensive final policy (74 SELL days) is the plausible binding constraint. Internal bull stances can remain buy-leaning while final signals stay defensive.

0.4.4 Figures




Figure 1: Agent vs buy-and-hold equity (2×2 grid).




Figure 2: Total returns (four windows).




Figure 3: Signal counts.




Figure 4: Bullish analyst-outlook share.

0.5 Conclusion

We presented a reproducible point-in-time LangGraph stack and a **scoped four-window case study**. Bear windows show lower loss and drawdown than B&H; bull outcomes are heterogeneous—an **architectural boundary** of defensive coordination, not a universal alpha claim.

Limitations: N=1 per window; two tickers; one LLM; anonymization not ablated; no multi-seed or index-universe evaluation; no latency CBS. **Future work:** seeds, broader PIT universes, anonymization ablations.

0.6 Contributions

1. Reproducible backtest harness and frozen CSV artifacts.
2. Point-in-time policy and optional anonymization.
3. Structured Pydantic stage outputs.
4. Multi-agent graph with tool-bound analyst decomposition.
5. Honest regime-conditioned case study with documented limits.

Acknowledgements

The agent graph implementation builds on **TradingAgents** (Xiao et al., 2024; Tauric Research, Apache License 2.0).

References

- Bădică, C., Bădică, A., Ganzha, M., Ivanovi'c, M., Paprzycki, M., Selişteanu, D. and Wrona, Z., 2025. Contemporary Agent Technology: LLM-Driven Advancements vs Classic Multi-Agent Systems. *arXiv preprint* arXiv:2509.02515 [cs.MA]. <https://arxiv.org/abs/2509.02515>
- Bailey, D.H., Borwein, J.M., L'opez de Prado, M. and Zhu, Q.J., 2014. Pseudo-Mathematics and Financial Charlatanism: The Effects of Backtest Overfitting on Out-of-Sample Perfor-

mance. *Notices of the American Mathematical Society*, 61(5), pp.458–471. <https://doi.org/10.1090/noti1105>

Dewansh, S., 2026. Policy Convergence and Explainable Stability in Orchestrated Multi-Agent LLM Trading Frameworks. *International Journal of Advanced Engineering and Management Science*, 12(2), pp.058–064. <https://doi.org/10.22161/ijaems.122.8>

Jeon, J. and Lee, H., 2026. Can Blindfolded LLMs Still Trade? An Anonymization-First Framework for Portfolio Optimization. *arXiv preprint arXiv:2603.17692* [cs.LG]. <https://arxiv.org/abs/2603.17692>

Li, X., Zeng, Y., Xing, X., Xu, J. and Xu, X., 2025. QuantAgents: Towards Multi-agent Financial System via Simulated Trading. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp.17438–17464, Suzhou, China. <https://doi.org/10.18653/v1/2025.findings-emnlp.945>

Miyazaki, K., Kawahara, T., Roberts, S. and Zohren, S., 2026. Toward Expert Investment Teams: A Multi-Agent LLM System with Fine-Grained Trading Tasks. *arXiv preprint arXiv:2602.23330* [cs.AI]. <https://arxiv.org/abs/2602.23330>

Nguyen, P. and Pham, T., 2026. Toward Reliable Evaluation of LLM-Based Financial Multi-Agent Systems: Taxonomy, Coordination Primacy, and Cost Awareness. *arXiv preprint arXiv:2603.27539* [cs.MA]. <https://arxiv.org/abs/2603.27539>

Papadakis, C., Filandrianos, G., Dimitriou, A., Lymperaiou, M., Thomas, K. and Stamou, G., 2025. StockSim: A Dual-Mode Order-Level Simulator for Evaluating Multi-Agent LLMs in Financial Markets. *arXiv preprint arXiv:2507.09255* [cs.CE]. <https://arxiv.org/abs/2507.09255>

Qian, L., Peng, X., Wang, Y., Zhang, V.J., He, H., Li, Z., Smith, H., Han, Y., He, Y., Li, H., Cao, Y., Yu, Y., Lopez-Lira, A., Lu, P., Nie, J.-Y., Xiong, G., Huang, J. and Ananiadou, S., 2026. When Agents Trade: Live Multi-Market Trading Arena for LLM Agents. In

Proceedings of the ACM Web Conference 2026. <https://doi.org/10.1145/3774904.3792821>

Xiao, Y., Sun, E., Luo, D. and Wang, W., 2024. TradingAgents: Multi-Agents LLM Financial Trading Framework. *arXiv preprint* arXiv:2412.20138 [q-fin.TR]. <https://arxiv.org/abs/2412.20138>

Yang, H., Liu, X.-Y. and Wang, C.D., 2023. FinGPT: Open-Source Financial Large Language Models. *arXiv preprint* arXiv:2306.06031 [q-fin.ST]. <https://arxiv.org/abs/2306.06031>

Yu, Y., Li, H., Chen, Z., Jiang, Y., Li, Y., Zhang, D., Liu, R., Suchow, J.W. and Khashanah, K., 2023. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. *arXiv preprint* arXiv:2311.13743 [q-fin.CP]. <https://arxiv.org/abs/2311.13743>

Zhao, T., Lyu, J., Jones, S., Garber, H., Pasquali, S. and Mehta, D., 2025. AlphaAgents: Large Language Model based Multi-Agents for Equity Portfolio Constructions. *arXiv preprint* arXiv:2508.11152 [q-fin.ST]. <https://arxiv.org/abs/2508.11152>